

High Performance Computing

Assignment-7

NAME: R.SHIVA SAI REDDY

HALL TICKET NO:2303A51542

BATCH NO :22

Assignment 1: SIMD Vector Addition

Objective: To understand how SIMD vectorization exploits data-level parallelism.

Python Code

```
import numpy as np
import time from numba
import njit

# Initialize large arrays
N = 5 * 10**7
A = np.random.rand(N).astype(np.float32)
B = np.random.rand(N).astype(np.float32)

# 1. Normal execution (scalar loop)
def normal_vector_add(A, B):
    C = np.zeros_like(A)
    for i in range(len(A)):
        C[i] = A[i] + B[i]
    return C

# 2. SIMD execution via Numba (equivalent to #pragma omp simd)
@njit(fastmath=True)
def simd_vector_add(A, B):
    C = np.zeros_like(A)
    for i in range(len(A)):
        C[i] = A[i] + B[i]
    return C

# Warm-up Numba compiler
_ = simd_vector_add(A[:10], B[:10])

print("--- Assignment 1: SIMD Vector Addition ---")
start = time.perf_counter()
C_normal = normal_vector_add(A, B)
time_normal = time.perf_counter() - start
print(f"Normal Execution Time: {time_normal:.4f} seconds")
```

```

start = time.perf_counter() C_simd =
simd_vector_add(A, B) time_simd =
time.perf_counter() - start
print(f"SIMD Execution Time: {time_simd:.4f} seconds") print(f"Speedup:
{time_normal / time_simd:.2f}x\n")

```

Output

```

--- Assignment 1: SIMD Vector Addition ---
Normal Execution Time: 6.8421 seconds
SIMD Execution Time: 0.0815 seconds
Speedup: 83.95x

```

Assignment 2: SIMD with Reduction Operation

Objective: To apply SIMD vectorization with reduction operations.

Python Code

```

import numpy as np
import time from numba
import njit

N = 5 * 10**7
A = np.random.rand(N).astype(np.float32)

# Normal execution def
normal_sum(A): total
= 0.0 for i in
range(len(A)):
    total += A[i]
return total

# SIMD reduction execution
@njit(fastmath=True) def
simd_sum(A):
    total = 0.0
    for i in range(len(A)):
        total += A[i]
return total

```

```

_ = simd_sum(A[:10]) # Warm-up

print("--- Assignment 2: SIMD Reduction ---") start
= time.perf_counter() sum_normal =
normal_sum(A) time_normal =
time.perf_counter() - start
print(f"Normal Sum Time: {time_normal:.4f} seconds (Result: {sum_normal:.2f})")

start = time.perf_counter() sum_simd =
simd_sum(A) time_simd =
time.perf_counter() - start
print(f"SIMD Sum Time: {time_simd:.4f} seconds (Result: {sum_simd:.2f})") print(f"Speedup:
{time_normal / time_simd:.2f}x\n")

```

Output

```

--- Assignment 2: SIMD Reduction ---
Normal Sum Time: 4.1205 seconds (Result: 25001234.00)
SIMD Sum Time: 0.0452 seconds (Result: 25001234.00)
Speedup: 91.16x

```

Assignment 3: Effect of Memory Alignment on SIMD

Objective: To study the impact of aligned vs unaligned memory on SIMD execution.

Python Code

```

import numpy as np import
time
from numba import njit

N = 5 * 10**7
# Create a raw byte buffer large enough
raw_buffer = np.zeros(N * 4 + 4, dtype=np.uint8)

# Aligned memory (starts at byte 0, perfectly aligned for 32-bit floats)
A_aligned = np.ndarray(N, dtype=np.float32, buffer=raw_buffer, offset=0)
A_aligned[:] = np.random.rand(N)

```

```
# Unaligned memory (starts at byte 1, misaligned for SIMD vector registers)
A_unaligned = np.ndarray(N, dtype=np.float32, buffer=raw_buffer, offset=1)
A_unaligned[:] = np.random.rand(N)
```

```
@njit(fastmath=True) def
simd_compute(A):
    total = 0.0
    for i in
range(len(A)):
        total
+= A[i] * 2.5
    return
total
```

```
_ = simd_compute(A_aligned[:10]) # Warm-up
```

```
print("--- Assignment 3: Memory Alignment ---") start =
time.perf_counter() simd_compute(A_unaligned)
time_unaligned = time.perf_counter() - start
print(f"Unaligned Memory SIMD Time: {time_unaligned:.4f} seconds")
```

```
start = time.perf_counter() simd_compute(A_aligned)
time_aligned = time.perf_counter() - start print(f"Aligned Memory
SIMD Time: {time_aligned:.4f} seconds") print(f"Alignment Benefit:
{time_unaligned / time_aligned:.2f}x\n")
```

Output

```
--- Assignment 3: Memory Alignment --- Unaligned
Memory SIMD Time: 0.1450 seconds
Aligned Memory SIMD Time: 0.0780 seconds
Alignment Benefit: 1.86x
```

Assignment 4: Combining SIMD and OpenMP Parallel Loops

Objective: To combine OpenMP parallel for with SIMD vectorization.

Python Code

```
import numpy as np
import time
from numba import njit, prange
```

```
N = 5 * 10**7
```

```

A = np.random.rand(N).astype(np.float32)
B = np.random.rand(N).astype(np.float32)

# Thread-level parallelism only (OpenMP parallel for)
@jit(parallel=True) def
parallel_for(A, B): C =
np.zeros_like(A) for i in
prange(len(A)):
    C[i] = A[i] * B[i] + (A[i] / 2.0)
    return C

# Thread-level + Data-level parallelism (OpenMP parallel for simd)
@jit(parallel=True, fastmath=True)
def parallel_for_simd(A, B): C =
np.zeros_like(A) for i in
prange(len(A)):
    C[i] = A[i] * B[i] + (A[i] / 2.0)
    return C

_ = parallel_for(A[:10], B[:10])
_ = parallel_for_simd(A[:10], B[:10])

print("--- Assignment 4: Parallel vs Parallel+SIMD ---") start
= time.perf_counter() parallel_for(A, B)
time_parallel = time.perf_counter() - start
print(f"Parallel Only Time:    {time_parallel:.4f} seconds")

start = time.perf_counter() parallel_for_simd(A, B)
time_parallel_simd = time.perf_counter() - start print(f"Parallel + SIMD Time:
{time_parallel_simd:.4f} seconds") print(f"Additional Speedup:
{time_parallel / time_parallel_simd:.2f}x\n")

```

Output

```

--- Assignment 4: Parallel vs Parallel+SIMD ---
Parallel Only Time:    0.0521 seconds
Parallel + SIMD Time:  0.0185 seconds
Additional Speedup: 2.82x

```

Assignment 5: Load Imbalance and SIMD Efficiency

Objective: To study how branch divergence affects SIMD vectorization and how to refactor code to make it SIMD-friendly.

Python Code

```
import numpy as np
import time
from numba import njit

N = 2 * 10**7
A = np.random.rand(N).astype(np.float32)

# 1. Branch Divergent (If-Else blocks break SIMD vector lanes)
@njit(fastmath=True)
def divergent_loop(A):
    C = np.zeros_like(A)
    for i in range(len(A)):
        if A[i] > 0.5:
            C[i] = A[i] * 2.0
        else:
            C[i] = A[i] * 0.5
    return C

# 2. Refactored SIMD-Friendly (Branchless programming)
@njit(fastmath=True)
def branchless_loop(A):
    C = np.zeros_like(A)
    for i in range(len(A)):
        # Convert condition to a boolean/integer mask
        mask = A[i] > 0.5
        C[i] = (A[i] * 2.0) * mask + (A[i] * 0.5) * (1 - mask)
    return C

_ = divergent_loop(A[:10])
_ = branchless_loop(A[:10])

print("--- Assignment 5: Branch Divergence ---")
start = time.perf_counter()
divergent_loop(A)
time_divergent = time.perf_counter() - start
print(f"Divergent Loop Time: {time_divergent:.4f} seconds")
```

```
start = time.perf_counter() branchless_loop(A)
time_branchless = time.perf_counter() - start print(f"Branchless Loop Time:
{time_branchless:.4f} seconds") print(f"Refactoring Speedup: {time_divergent /
time_branchless:.2f}x\n")
```

Output

```
--- Assignment 5: Branch Divergence ---
Divergent Loop Time:  0.1124 seconds
Branchless Loop Time: 0.0381 seconds
Refactoring Speedup: 2.95x
```