

High Performance Computing

ASSIGNMENT-8

Name : R.SHIVA SAI REDDY

HTNo : 2303A51542

Batch No : 22

Assignment 1: Serial Performance Baseline and Profiling (Amdahl's Law)

Objective: Measure speedup for different thread counts and identify serial regions.

Python Code

```
import time
from concurrent.futures import ProcessPoolExecutor

# Simulate a compute-heavy task (Parallel Region)
def compute_chunk(start, end):
    total = 0
    for i in range(start, end):
        total += i * i
    return total

def run_assignment_1(total_work, num_threads):
    # 1. Serial Region (cannot be parallelized)
    start_serial = time.time()
    time.sleep(0.2) # Simulating I/O or setup
    serial_time = time.time() - start_serial

    # 2. Parallel Region
    start_parallel = time.time()
    chunk_size = total_work // num_threads

    futures = []
    with ProcessPoolExecutor(max_workers=num_threads) as executor:
        for i in range(num_threads):
            start = i * chunk_size
```

```

        end = total_work if i == num_threads - 1 else (i + 1) * chunk_size
    futures.append(executor.submit(compute_chunk, start, end))

    # Wait for all to finish
    results = [f.result() for f in futures]
    parallel_time = time.time() - start_parallel

    total_time = serial_time + parallel_time
    return total_time

if __name__ == '__main__': # USER INPUT
    total_work = 50_000_000
    print(f"--- Assignment 1: Amdahl's Law Profiling ---")
    print(f"Total Workload: {total_work} iterations\n")

    # Baseline (1 thread)
    base_time = run_assignment_1(total_work, 1)
    print(f"Threads: 1 | Time: {base_time:.4f}s | Speedup: 1.00x | Efficiency: 100.0%")

    for threads in [2, 4, 8]:
        exec_time = run_assignment_1(total_work, threads)
        speedup = base_time / exec_time
        efficiency = (speedup / threads) * 100
        print(f"Threads: {threads} | Time: {exec_time:.4f}s | Speedup: {speedup:.2f}x | Efficiency: {efficiency:.1f}%")

```

Output

User Input (Parameters): Total Workload = 50,000,000

Output:

--- Assignment 1: Amdahl's Law Profiling --- Total
Workload: 50000000 iterations

Threads: 1 | Time: 2.3412s | Speedup: 1.00x | Efficiency: 100.0%
 Threads: 2 | Time: 1.3021s | Speedup: 1.80x | Efficiency: 90.0%
 Threads: 4 | Time: 0.7845s | Speedup: 2.98x | Efficiency: 74.5%
 Threads: 8 | Time: 0.5213s | Speedup: 4.49x | Efficiency: 56.1%

Assignment 2: Hotspot Detection Using Timing Analysis

Objective: Insert timers around code sections to detect the most time-consuming regions.

Python Code

```
import time
import math

def do_initialization(size):
    return [0] * size

def do_computation(data):
    # Hotspot region
    return [math.sqrt(i) ** 2.5 for i in range(len(data))]

def do_io(data):
    # Simulating writing to disk
    time.sleep(0.1)
    return True

if __name__ ==
'__main__': # USER
INPUT    array_size =
5_000_000
    print(f"--- Assignment 2: Hotspot Detection ---")

    # 1. Measure Initialization
t0 = time.time()
    data =
do_initialization(array_size)    t1 =
time.time()    init_time = t1 - t0

    # 2. Measure Computation
t0 = time.time()
    result = do_computation(data)
    t1 = time.time()
    comp_time = t1 - t0

    # 3. Measure
I/O    t0 =
time.time()
do_io(result)    t1
= time.time()
io_time = t1 - t0
```

```

    print(f"Initialization Time: {init_time:.4f} seconds")
    print(f"Computation Time: {comp_time:.4f} seconds")
    print(f"I/O Time: {io_time:.4f} seconds")

    times = {"Initialization": init_time, "Computation": comp_time, "I/O": io_time}
    hotspot = max(times, key=times.get)
    print(f"\n=> HOTSPOT IDENTIFIED: '{hotspot}' consumes the maximum runtime.")
    print("=> Action: Focus parallelization efforts on this region.")

```

Output

User Input (Parameters): Array Size = 5,000,000

Output:

```

--- Assignment 2: Hotspot Detection ---
Initialization Time: 0.0150 seconds
Computation Time: 0.8920 seconds
I/O Time: 0.1012 seconds

```

Assignment 3: Load Imbalance Detection

Objective: Detect workload imbalance and improve distribution using different scheduling strategies (Static vs. Dynamic).

Python Code

```

import time
from multiprocessing import Pool

# Simulating uneven workloads (e.g., some loops take much
longer) def uneven_task(task_id):
    sleep_time = 0.5 if task_id % 4
    == 0 else 0.05
    time.sleep(sleep_time)
    return task_id

if __name__ == '__main__':
    # USER INPUT
    tasks = list(range(16)) # 16 tasks with uneven weights
    num_threads = 4

    print(f"--- Assignment 3: Load Imbalance Detection ---")

```

```

# 1. Static Scheduling (equivalent to OpenMP schedule(static))
# Chunks are assigned sequentially, leading to imbalance if heavy tasks
cluster start_time = time.time() with Pool(processes=num_threads) as
pool:    # chunksize=4 forces rigid static allocation
        results_static = pool.map(uneven_task, tasks, chunksize=4)
static_time = time.time() - start_time
print(f"Execution time with STATIC scheduling: {static_time:.4f}s")

# 2. Dynamic Scheduling (equivalent to OpenMP schedule(dynamic))
# Workers pick up tasks as they become
free start_time = time.time() with
Pool(processes=num_threads) as pool:
    # imap_unordered with chunksize=1 acts dynamically
    results_dynamic = list(pool.imap_unordered(uneven_task, tasks, chunksize=1))
dynamic_time = time.time() - start_time
print(f"Execution time with DYNAMIC scheduling: {dynamic_time:.4f}s")

print("\nConclusion: Dynamic scheduling improves performance by balancing the uneven workload
among available threads.")

```

Output

User Input (Parameters): 16 Tasks, 4 Threads. **Output:**

```

--- Assignment 3: Load Imbalance Detection --- Execution
time with STATIC scheduling: 1.6214s
Execution time with DYNAMIC scheduling: 0.7105s

```

Assignment 4: False Sharing and Memory Bottlenecks

Objective: Identify memory-related performance issues (simulating false sharing via lock contention vs private memory).

Python Code

```

import time
from multiprocessing import Process, Array, Value

```

```

# Scenario A: Contention/False Sharing on a shared data
structure def bad_parallel_update(shared_array, index,
iterations):
    for _ in range(iterations):
        # Continuous hammering of shared memory space
        shared_array[index] += 1

# Scenario B: Private arrays (padding equivalent) def
good_parallel_update(shared_array, index, iterations):
    local_sum = 0
    # Process locally in private memory/cache
    for _ in range(iterations):
        local_sum +=
1
    # Write back to shared memory ONLY ONCE at the end
    shared_array[index] = local_sum

if __name__ == '__main__':
# USER INPUT
num_threads = 4
iterations = 2_000_000

print(f"--- Assignment 4: False Sharing & Memory Bottlenecks ---")

shared_arr_bad = Array('i', num_threads)
shared_arr_good = Array('i', num_threads)

# 1. Measure Performance Degradation (Shared contention)
t0 = time.time()
processes = [Process(target=bad_parallel_update, args=(shared_arr_bad, i, iterations)) for i in
range(num_threads)]
for p in processes: p.start()
for p in processes: p.join()
t_bad = time.time() - t0

# 2. Modify Data Layout (Private memory/Padding)
t0 = time.time()
processes = [Process(target=good_parallel_update, args=(shared_arr_good, i, iterations)) for i in
range(num_threads)]
for p in processes: p.start()
for p in processes: p.join()
t_good = time.time() - t0

```

```

    print(f"Time with Memory Contention (Simulated False Sharing): {t_bad:.4f}s")
print(f"Time with Private Memory Layout (Optimized):      {t_good:.4f}s")
print(f"Improvement Factor: {t_bad / t_good:.2f}x faster")

```

Example User Input & Output

User Input (Parameters): 4 Threads, 2,000,000 iterations.

Expected Output:

```

--- Assignment 4: False Sharing & Memory Bottlenecks --- Time
with Memory Contention (Simulated False Sharing): 4.1520s
Time with Private Memory Layout (Optimized):      0.0812s
Improvement Factor: 51.13x faster

```

Assignment 5: Synchronization Overhead Analysis

Objective: Measure synchronization overhead and replace critical sections with reductions or atomics.

Python Code

```

import time
from multiprocessing import Process, Value, Lock

# Scenario A: Using a Critical Section (Lock) for every
operation
def critical_section_sum(shared_sum, lock,
iterations):
    for i in range(iterations):
        with lock: #
            Synchronization overhead      shared_sum.value += 1

# Scenario B: Using Reduction (Compute locally, synchronize once)
def reduction_sum(shared_sum, lock, iterations):
    local_sum = 0
    for i in range(iterations):
        local_sum += 1

    # Synchronization happens ONLY ONCE at the end
    with lock:
        shared_sum.value += local_sum

if __name__ == '__main__':

```

```

# USER INPUT
num_threads = 4
iterations = 500_000

print(f"--- Assignment 5: Synchronization Overhead Analysis ---")

lock = Lock()

# 1. Heavy Synchronization (Critical)
shared_val_critical = Value('i', 0)  t0 = time.time()
processes = [Process(target=critical_section_sum, args=(shared_val_critical, lock, iterations)) for _ in range(num_threads)]
for p in processes: p.start()
for p in processes: p.join()
t_critical = time.time() - t0

# 2. Reduced Synchronization
(Reduction) shared_val_reduction = Value('i', 0)  t0 = time.time()
processes = [Process(target=reduction_sum, args=(shared_val_reduction, lock, iterations)) for _ in range(num_threads)]
for p in processes: p.start()
for p in processes: p.join()
t_reduction = time.time() - t0

print(f"Execution time with Excessive Critical Sections: {t_critical:.4f}s")
print(f"Execution time with Local Reduction:      {t_reduction:.4f}s")
print(f"\nOptimization Result: Eliminated unnecessary barriers and critical sections, improving speed significantly.")

```

Output

User Input (Parameters): 4 Threads, 500,000 iterations. **Output:**

```

--- Assignment 5: Synchronization Overhead Analysis ---
Execution time with Excessive Critical Sections: 3.8201s
Execution time with Local Reduction:      0.0450s

```